

Data Management With R: Relational Data

Matthias Haber

16 October 2017

Last weeks Homework

Last weeks Homework

Prerequisites

Packages

```
library(tidyverse)  
library(readr)
```

```
# base GitHub url
url <- paste0("https://raw.githubusercontent.com/",
              "mhaber/HertieDataScience/master/",
              "slides/week6/data/")

# link to datasets
pew <- read_csv(paste0(url, "pew.csv"))
billboard <- read_csv(paste0(url, "billboard.csv"))
weather <- read_tsv(paste0(url, "weather.txt"))
```

Terminology

What is a database:

A collection of information organized to afford efficient retrieval.

“When people use the word database, fundamentally what they are saying is that the data should be self-describing and it should have a schema.” (Jim Gray)

Advantages of databases:

1. Sharing

- Support concurrent access by multiple readers and writers

2. Data Model Enforcement

- Make sure all applications see clean, organized data

3. Scale

- Work with datasets too large to fit into memory


4. Flexibility

- Use the data in new, unanticipated ways

Evaluate databases:


- How are the data physically organized on disk?
- What kinds of queries are efficiently supported, and which are not?
- How hard is it to update the data, or add new data
- What happens when I encounter new queries that I didn't anticipate. Do I reorganize the data?

Historical example: network databases



- All programs must be rewritten to accommodate reorganizations of the data
- difficult to look up all orders that involve a specific part

Historical example: hierarchical databases



- hierarchical order of data in terms of segments (e.g. IBM's IMS system that still exists)
 - works great if you want to find all orders for a particular customer
 - but what if you want to find all customers who ordered a Nail?

Key idea of relational databases

“Relational Database Management Systems were invented to let you use one set of data in multiple ways. Including ways that were unforeseen at the time the database is built, and then at the time the first applications are written.”
(Curt Monash)

God made the integers; all else is the work of man. (Leopold Kronecker, 19th century mathematician)

Cood made relations; all else is the work of man. (Raghu Ramakrishnan, DB text book author)

Relational databases concept

Relational databases (Edgar Codd 1970)

- Everything is a table
- Every row in a table has the same (types of) columns
- Relationships are implicit: no pointers

course	Student_ID
GRAD-E123	223
GRAD-E123	224
GRAD-E456	255
GRAD-E456	244

Student_ID	Student_Name
223	Jane
224	Joe

Relational databases history

- Pre relational: if your data changed, your application broke
- Early RDBMS were buggy and slow, but required only 5% of the application code

Notion of data independence:


“Activities of users at terminals and most application programs should remain unaffected when the internal representation of data is changed and even when some aspects of the external representation are changed.” (Ted Cod)

Relational algebra:


- Programs that manipulate a tabular data exhibit an **algebraic structure** that we can use to reason about them and manipulate the logical model independent of any physical data representation

Physical data independence

Access data through a high-level language (SQL), not physical pointers.



physical data independence



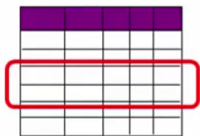
files and
pointers

```
SELECT seq
FROM ncbi_sequences
WHERE seq = 'GATTACGATATTA';
```

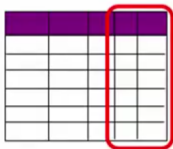
```
f = fopen('table_file');
fseek(10030440);
while (True) {
    fread(&buf, 1, 8192, f);
    if (buf == GATTACGATATTA) {
        . . .
```

Algebra of tables

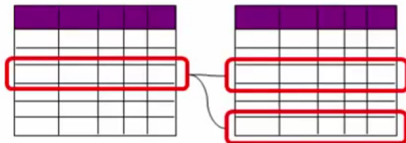
Use operations from relational algebra and set theory.



select



project



join

Other operators: aggregate, union, difference, cross product

$$N = ((z * 2) + ((z * 3) + 0))/1$$

Algebraic optimization

$$N = ((z * 2) + ((z * 3) + 0))/1$$

Algebraic laws:


1. $x + 0 = x$
2. $x/1 = x$
3. $(n * x + n * y) = n * (x + y)$
4. $x * y = y * x$

Apply rules 1,3,4,2: $N = (2 + 3) * z$

All relational databases do algebraic optimization when you write a query.

Relational algebra operators

- Union \cup , intersection \cap , difference $-$
- Selection σ
- Projection Π
- Join \bowtie
- Duplicate elimination ρ
- Grouping and aggregation g
- Sorting τ



Sets vs. bags

Sets: $\{a,b,c\}$, $\{a,d,e,f\}$, $\{\}$, ... Bags: $\{a,a,b,c\}$, $\{b,b,b,b,b\}$, ...

Relational algebra has two semantics:

- Set semantics = standard relational algebra
- Bag semantics = extended relational algebra

Rule of thumb:

- Every commercial database will assume bag semantics ##
Union

$R1 \cup R2$

```
SELECT * FROM R1
UNION
SELECT * FROM R2
```

Difference

$R1 - R2$

```
SELECT * FROM R1  
EXCEPT  
SELECT * FROM R2
```

R1			R2			R1 - R2	
A	B		A	B		A	B
a1	b1	-	a1	b1	=	a2	b1
a2	b1		a3	b4			

- Derived operator using minus:

$$R1 \cap R2 = R1 - (R1 - R2)$$

- Derived using joins:

$$R1 \cap R2 = R1 \bowtie R2$$

- Returns all tuples which satisfy a condition

$\sigma_c(R)$

- The condition c can be $=, <, \leq, >, \geq, <>$
- And you can also use boolean expression (and, or ,not)

Selection

Employee

SSN	Name	Salary
1234545	John	20000
5423341	Smith	60000
4352342	Fred	50000

$\sigma_{\text{Salary} > 40000}$ (Employee)

SSN	Name	Salary
5423341	Smith	60000
4352342	Fred	50000

- Eliminates columns

$$\Pi_{A_1, \dots, A_n}(R)$$

Projection

Employee

SSN	Name	Salary
1234545	John	20000
5423341	John	60000
4352342	John	20000

$\Pi_{\text{Name, Salary}}(\text{Employee})$

Name	Salary
John	20000
John	60000
John	20000

Name	Salary
John	20000
John	60000

•

Set semantics

Cross Product

- Each tuple in $R1$ with each tuple in $R2$: $R1 \times R2$
- Rare in practice, but can come up in analytics
 - e.g. find all pairs of similar tweets

Cross Product

Employee

Name	SSN
John	999999999
Tony	777777777

Dependent


EmpSSN	DepName
999999999	Emily
777777777	Joe

Employee \bowtie Dependent

Name	SSN	EmpSSN	DepName
John	999999999	999999999	Emily
John	999999999	777777777	Joe
Tony	777777777	999999999	Emily
Tony	777777777	777777777	Joe


Join

A join is a way of connecting each row in x to zero, one, or more rows in y . The following diagram shows each potential match as an intersection of a pair of lines.



Inner join

The simplest type of join is the **inner join**. An inner join matches pairs of observations whenever their keys are equal. Unmatched rows are not included in the result



Outer joins


An inner join keeps observations that appear in both tables. An **outer join** keeps observations that appear in at least one of the tables. There are three types of outer joins:

- A **left join** keeps all observations in *x*.
- A **right join** keeps all observations in *y*.
- A **full join** keeps all observations in *x* and *y*.

Left



Right



SQL

Structure Query Language is a language for interacting with databases

- A query is a request for data from a database table (or combination of tables)
- SQL can be used to query but also to create and modify databases.

SELECT

In SQL, you can select data from a table using a `SELECT` statement. For example, the following query selects the `name` column `FROM` the `people` table:

```
SELECT name  
FROM people
```

`SELECT` and `FROM` are keywords. They are not case-sensitive but it's good practice to make SQL keywords uppercase to distinguish them from other parts of your query, like column and table names.

SELECT

To select multiple columns from a table, simply separate the column names with commas!

```
SELECT name, birthdate  
FROM people
```

To select all columns

```
SELECT *  
FROM people
```

You can also LIMIT the number of rows returned:

```
SELECT *  
FROM people  
LIMIT 10
```

SELECT DISTINCT

You can use the DISTINCT keyword to select all the unique values from a column.

```
SELECT DISTINCT language  
FROM films
```

The COUNT statement returns the number of rows in one or more columns.

```
SELECT COUNT(*)  
FROM people
```

It's common to combine COUNT with DISTINCT to count the number of distinct values in a column.

```
SELECT COUNT(DISTINCT birthdate)  
FROM people
```

WHERE

The WHERE keyword allows you to filter based on both text and numeric values in a table. The WHERE clause always comes after the FROM statement.

```
SELECT title  
FROM films  
WHERE title = 'Metropolis'
```


You can build up your WHERE queries by combining multiple conditions with the AND keyword.

```
SELECT title  
FROM films  
WHERE release_year > 1994  
AND release_year < 2000
```

WHERE AND OR

The OR operator allows you to select rows based on multiple conditions where some but not all of the conditions need to be met?

```
SELECT title  
FROM films  
WHERE release_year = 1994  
OR release_year = 2000
```

BETWEEN

The BETWEEN keyword provides a useful shorthand for filtering values within a specified range. BETWEEN is always inclusive can be used with multiple AND and OR operators

```
SELECT title  
FROM films  
WHERE release_year  
BETWEEN 1994 AND 2000
```

WHERE IN

The IN operator allows you to specify multiple values in a WHERE clause, making it easier and quicker to specify multiple OR conditions.

```
SELECT name  
FROM kids  
WHERE age IN (2, 4, 6, 8, 10)
```

NULL and IS NULL

NULL represents a missing or unknown value. You can check for NULL values using the expression `IS NULL`.

```
SELECT COUNT(*)  
FROM people  
WHERE birthdate IS NULL
```

You can use the `IS NOT NULL` operator to filter out missing values.

LIKE and NOT LIKE

The LIKE operator can be used in a WHERE clause to search for a pattern in a column. You use a wildcard as a placeholder for some other values. There are two wildcards you can use with LIKE:

- The % wildcard will match zero, one, or many characters in text:
- The _ wildcard will match a single character.

For example,

```
SELECT name  
FROM companies  
WHERE name LIKE 'Data%'
```

finds 'Data', 'Datab', 'Database'.

Aggregate functions

SQL provides a few functions to perform some calculation on the data in a database. Those aggregate functions can also be combined with the WHERE clause

For example,

```
SELECT AVG(budget)
FROM films
```

gives you the average value from the budget column of the films table. Similarly, the MAX function returns the highest budget:

```
SELECT MAX(budget)
FROM films
```

The SUM function returns the result of adding up the numeric values in a column:

Arithmetics

You can perform basic arithmetic with symbols like +, -, *, and /.

For example, this gives a result of 12:

```
SELECT (4 * 3)
```

However, the following gives a result of 1,

```
SELECT (4 / 3)
```

SQL assumes that if you divide an integer by an integer, you want to get an integer back. If you want more precision when dividing, you can add decimal places to your numbers.

You use the AS keyword to assign a temporary name to something (i.e. alias).

```
SELECT MAX(budget) AS max_budget,  
       MAX(duration) AS max_duration  
FROM films
```

ORDER BY

The `ORDER BY` keyword is used to sort results according to the values of one or more columns. By default `ORDER BY` will sort in ascending order. If you want to sort the results in descending order, you can use the `DESC` keyword.

```
SELECT title  
FROM films  
ORDER BY release_year DESC
```

ORDER BY multiple columns

ORDER BY can also be used to sort on multiple columns. It will sort by the first column specified, then sort by the next, then the next, and so on.

```
SELECT birthdate, name  
FROM people  
ORDER BY birthdate, name
```

GROUP BY

GROUP BY allows you to group a result by one or more columns. GROUP BY always goes after the FROM clause.

```
SELECT sex, count(*)  
FROM employees  
GROUP BY sex
```

Commonly, GROUP BY is used with aggregate functions like COUNT() or MAX().

In SQL, aggregate functions can't be used in WHERE clauses. If you want to filter based on the result of an aggregate function, you need to use the HAVING clause.

```
SELECT release_year  
FROM films  
GROUP BY release_year  
HAVING COUNT(title) > 10
```

INNER JOIN

INNER JOIN matches pairs of observations whenever their keys are equal.

```
SELECT *  
FROM left_table  
INNER JOIN right_table  
ON left_table.id = right_table.id
```

User-defined function support

- PostgreSQL
 - SQL, PL/pSQL, Python, C/C++, R
- Microsoft SQL Server
 - SQL, T-SQL, C#
- Oracle
 - SQL, PL-SQL, Java, C/C++, Python, others
- SQLite
 - None

Homework Exercises

Homework Exercises

For this week's homework exercises go to Moodle and answer the Quiz posted in the Tidy Data section.

Deadline: Sunday, October 15 before midnight.

That's it for today. Questions?